

hw2


January 23, 2025

```
[12]: import numpy as np
import scipy as sp
import matplotlib.pyplot as plt
```

1 1. Subset Selection and Simulations

1. Code:

```
[13]: N = 500
beta = np.array([[4],[2],[2],[2]])
rng = np.random.default_rng()
beta_hat = []
for i in range(500):
    X = rng.normal(0, 1, (N, 4))
    U = rng.normal(0, 2, (N, 1))
    Y = X @ beta + U
    R = 0
    ix = 0
    beta_i = 0
    for i in range(4):
        reg = sp.stats.linregress(X[:, i], Y[:, 0])
        if (reg.rvalue > R):
            R = reg.rvalue
            ix = i
            beta_i = reg.slope
    beta_hat.append([ix + 1, beta_i])
beta_hat = np.array(beta_hat)
plt.ylabel(r"$\hat{\beta}_{\hat{j}}$ and $\hat{\beta}_j$")
plt.xlabel("Simulations")
plt.plot(beta_hat[:, 0], label = "Index")
plt.plot(beta_hat[:, 1], label = "Estimate")
plt.legend()
plt.show()
```




2. The trick applied in class, namely utilizing the standardized OLS coefficients, and simply comparing magnitudes does work here. In this case, there is no deleterious impact of considering standardized OLS coefficients, since, by design, each of the x_i 's have the same variance. Additionally, since this is just a toy example, and we are simply trying to observe the relationships between the β_i 's, and not attempting causal inference on the features, we face no problem again. However, it is not super important, since the features being orthogonal by design already lends itself to standardized regression coefficients.
3. In every single simulation, the researcher is able to identify the most relevant regressor. This is not surprising, since, by design, the features are all orthogonal. As a result, we are very easily able to identify the most significant covariate.
4. Code:

```
[14]: beta_hat = []
gamma = 0.5
cov = np.array([[1, 0, 0, 0],
                [0, 1, gamma, gamma],
                [0, gamma, 1, gamma],
                [0, gamma, gamma, 1]])
for i in range(500):
    X = rng.multivariate_normal(np.zeros(4), cov=cov, size = (N, ))
```

```

U = rng.normal(0, 2, (N, 1))
Y = X @ beta + U
R = 0
j_hat = 0
beta_j = 0
for j in range(4):
    reg = sp.stats.linregress(X[:,j], Y[:, 0])
    if (reg.rvalue >= R):
        j_hat = j
        beta_j = reg.slope
        R = reg.rvalue
    beta_hat.append([j_hat + 1, beta_j])
beta_hat = np.array(beta_hat)
plt.ylabel(r"$\hat{\beta}_j$ and $\hat{j}$")
plt.xlabel("Simulations")
plt.plot(beta_hat[:, 0], label = "Index")
plt.plot(beta_hat[:,1], label = "Estimate")
plt.legend()
plt.show()
print(f"The proportion that the researcher correctly identifies is {np.
    ↳mean(beta_hat[:, 0] == np.full((N,), 1))}")

```



The proportion that the researcher correctly identifies is 0.302

5. The researcher is able to identify the most relevant regressor approximately $\frac{1}{3}$ of the time. Considering the features are very much not orthogonal now, this is not surprising.
6. This is directly related to omitted variable bias because there are covariates excluded from the model that have non-zero effect on the outcome, and are correlated with the covariates in the model.
7. Code:

```
[15]: def simulation(gamma):
    beta_hat = []
    cov = np.array([[1, 0, 0, 0],
                    [0, 1, gamma, gamma],
                    [0, gamma, 1, gamma],
                    [0, gamma, gamma, 1]])
    for _ in range(500):
        X = rng.multivariate_normal(np.zeros(4), cov=cov, size = (N, ))
        U = rng.normal(0, 2, (N, 1))
        Y = X @ beta + U
        R = 0
        j_hat = 0
        beta_j = 0
        for j in range(4):
            reg = sp.stats.linregress(X[:,j], Y[:, 0])
            if (reg.rvalue > R):
                j_hat = j
                beta_j = reg.slope
                R = reg.rvalue
        beta_hat.append([j_hat + 1, beta_j])
    beta_hat = np.array(beta_hat)
    return [gamma, np.mean(beta_hat[:, 0] == np.full((N,), 1))]
```


The following are the results of running a few choices of gamma:

```
[16]: res = []
    for gamma in np.linspace(-1, 1, 65):
        res.append(simulation(gamma))
    res = np.array(res)
    plt.plot(res[:, 0], res[:, 1])
    res = res[res[:, 1] > 0.95]
    gamma = np.max(res, axis= 0)
    plt.axhline(y = 0.95)
    plt.axvline(x = gamma[0])
    plt.xlabel(r"$\gamma$")
    plt.ylabel("Proportion Correctly Identified")
    plt.show()
```

C:\Users\matth\AppData\Local\Temp\ipykernel_1804\2683143143.py:8:

RuntimeWarning: covariance is not symmetric positive-semidefinite.

```
X = rng.multivariate_normal(np.zeros(4), cov=cov, size = (N, ))
```



```
[17]: print(gamma[0])
```

0.34375

For values of γ in $(0, 0.34375)$, we select the most significant β at least 0.95 of the time.

8. Code:

```
[18]: def simulation(gamma, beta_1, N, num_sim):
    beta = [[beta_1], [2], [2], [2]]
    beta_hat = []
    cov = np.array([[1, 0, 0, 0],
                    [0, 1, gamma, gamma],
                    [0, gamma, 1, gamma],
                    [0, gamma, gamma, 1]])
    for _ in range(num_sim):
        X = rng.multivariate_normal(np.zeros(4), cov=cov, size = (N, ))
        U = rng.normal(0, 2, (N, 1))
        Y = X @ beta + U
        R = 0
```

```


j_hat = 0
beta_j = 0
for j in range(4):
    reg = sp.stats.linregress(X[:,j], Y[:, 0])
    if (reg.rvalue > R):
        j_hat = j
        beta_j = reg.slope
        R = reg.rvalue
    beta_hat.append([j_hat + 1, beta_j])
beta_hat = np.array(beta_hat)
return [beta_1, gamma, np.mean(beta_hat[:, 0] == np.full((num_sim,), 1))]

```

```

[21]: res = []
for beta in np.linspace(1, 8, 65):
    res.append(simulation(1/2, beta, 500, 500))
res = np.array(res)
plt.plot(res[:, 0], res[:, 2])
beta = np.min(res[res[:, 2] > 0.95], axis= 0)[0]
plt.axhline(y = 0.95)
plt.axvline(x = beta)
plt.xlabel(r"$\beta$")
plt.ylabel("Proportion Correctly Identified")
plt.show()

```



```
[23]: print(f"The threshold of beta_1 is {beta}")
```

The threshold of beta_1 is 4.609375

Side Note: these curves look logistic. Here's some code to analyze this.

```
[24]: beta_res = []
for beta in np.linspace(1, 8, 65):
    sim_res = simulation(1/2, beta, 500, 500)
    beta_res.append([sim_res[0], sim_res[2]])
gamma_res = []
for gamma in np.linspace(-1/3, 1, 65):
    sim_res = simulation(gamma, 4, 500, 500)
    gamma_res.append([sim_res[1], sim_res[2]])
beta_res = np.array(beta_res)
gamma_res = np.array(gamma_res)
def logifunc(x,A,x0,k):
    return A / (1 + np.exp(-k*(x-x0)))
gamma_param, gamma_cov = sp.optimize.curve_fit(logifunc, gamma_res[:, 0],
↪gamma_res[:, 1])
beta_param, beta_cov = sp.optimize.curve_fit(logifunc, beta_res[:, 0],
↪beta_res[:, 1])
```

```
[25]: gamma_param, gamma_cov = sp.optimize.curve_fit(logifunc, gamma_res[:, 0],
↪gamma_res[:, 1])
beta_param, beta_cov = sp.optimize.curve_fit(logifunc, beta_res[:, 0],
↪beta_res[:, 1])
plt.figure()
plt.plot(gamma_res[:, 0], logifunc(gamma_res[:, 0], *gamma_param), label =
↪"predicted")
plt.plot(gamma_res[:, 0], gamma_res[:, 1], label = "True")
gamma_stderr = np.sqrt(np.diag(gamma_cov))
gamma_cis = np.vstack((gamma_param - 2.576 * gamma_stderr, gamma_param + 2.576
↪* gamma_stderr)).T
print(f"The best fit logistic function for gamma is: \n A = {gamma_param[0]},
↪x0 = {gamma_param[1]}, k = {gamma_param[2]}")
print(f"99% confidence intervals: \n{gamma_cis}")
plt.legend()
plt.title("Gamma vs predicted")
plt.figure()
plt.plot(beta_res[:, 0], logifunc(beta_res[:, 0], *beta_param), label =
↪"predicted")
plt.plot(beta_res[:, 0], beta_res[:, 1], label = "true")
plt.legend()
beta_stderr = np.sqrt(np.diag(beta_cov))
```

```

beta_cis = np.vstack((beta_param - 2.576 * beta_stderr, beta_param + 2.576 *
    ↪beta_stderr)).T
print(f"The best fit logistic function for beta is: \n A = {beta_param[0]}, x0_
    ↪= {beta_param[1]}, k = {beta_param[2]}")
print(f"99% confidence intervals: \n{beta_cis}")
plt.title("Beta vs predicted")
plt.show()

```

The best fit logistic function for gamma is:

A = 1.0014851104564708, x0 = 0.4678934046299783, k = -25.07689491085193

99% confidence intervals:

```

[[ 0.99835942  1.0046108 ]
 [ 0.46659315  0.46919365]
 [-25.79014223 -24.36364759]]

```

The best fit logistic function for beta is:

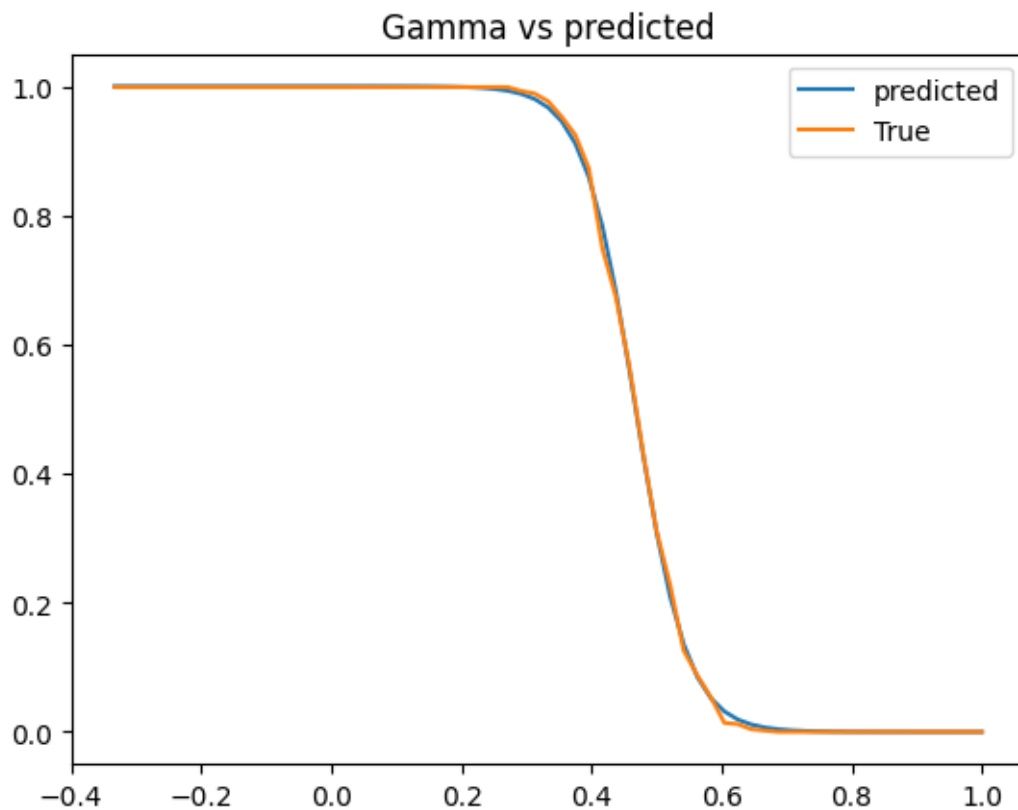
A = 1.0007931910542014, x0 = 4.1366860044946145, k = 6.302759627864724


99% confidence intervals:

```

[[0.99788885 1.00369753]
 [4.13133246 4.14203955]
 [6.11691148 6.48860778]]

```





Whether subset selection works or not seems to follow an asymptotically logistic distribution, for both β and γ .

2 2 Subset selection and data: growth regressions

1. Imports:

```
[26]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import scipy as sp
from sklearn.linear_model import LinearRegression
```

Read data

```
[27]: data = pd.read_excel("millions.xls", "BARROSHO", index_col= "code", na_values=".
↵")
labels = pd.read_excel("millions.XLS", "Variable code", index_col="#")
```

```
[28]: data.head()
```

```
[28]: #      country      gamma      X1      X2      X3      X4      X5      X6 \
code
DZA  1      Algeria  0.013690  7.438972  47.299999  0.46      0      0  0.131
AGO  2      Angola   0.000569  6.786717      NaN  0.21      1      0  0.281
BEN  3      Benin   -0.006586  7.019297  38.900002  0.27      1      0  0.050
BWA  4      Botswana 0.056195  6.284134  45.700001  0.42      1      0  0.072
HVO  5  Burkina Faso 0.004206  6.152733  36.299999  0.08      1      0  0.050

      X7  ...  X53  X54      X55      X56      X57      X58      X59  X60 \
code      ...
DZA  186.555466  ...  0.0  0.0  0.005  0.99  0.005  2855.520020  0.196  0.0
AGO  239.605362  ...  0.0  0.0  0.000  0.00  0.150  2319.385498  0.268  0.0
BEN   4.489583  ...  0.0  0.0  0.000  0.15  0.080  1372.623291  0.009  0.0
BWA  23.388475  ...  0.0  0.0  0.000  0.00  0.250  210.918488  0.533  5.0
HVO   4.489583  ...  0.0  0.0  0.000  0.25  0.000      NaN  0.001  1.0

      X61  X62
code
DZA   0.836  0.0
AGO   0.000  0.0
BEN   0.000  0.0
BWA   0.000  0.0
HVO   0.000  0.0
```

[5 rows x 64 columns]

```
[29]: data.drop(["#", "country"], axis = 1, inplace = True)
```

```
[30]: print(data.head())
print(labels.head())
```

```
      gamma      X1      X2      X3      X4      X5      X6      X7      X8 \
code
DZA  0.013690  7.438972  47.299999  0.46      0      0  0.131  186.555466  21.069
AGO  0.000569  6.786717      NaN  0.21      1      0  0.281  239.605362      NaN
BEN -0.006586  7.019297  38.900002  0.27      1      0  0.050   4.489583      NaN
BWA  0.056195  6.284134  45.700001  0.42      1      0  0.072  23.388475      NaN
HVO  0.004206  6.152733  36.299999  0.08      1      0  0.050   4.489583      NaN

      X9  ...  X53  X54      X55      X56      X57      X58      X59  X60 \
code      ...
DZA  13.303  ...  0.0  0.0  0.005  0.99  0.005  2855.520020  0.196  0.0
AGO   NaN  ...  0.0  0.0  0.000  0.00  0.150  2319.385498  0.268  0.0
BEN   NaN  ...  0.0  0.0  0.000  0.15  0.080  1372.623291  0.009  0.0
BWA   NaN  ...  0.0  0.0  0.000  0.00  0.250  210.918488  0.533  5.0
HVO   NaN  ...  0.0  0.0  0.000  0.25  0.000      NaN  0.001  1.0

      X61  X62
```

```
code
DZA    0.836  0.0
AGO    0.000  0.0
BEN    0.000  0.0
BWA    0.000  0.0
HVO    0.000  0.0
```

```
[5 rows x 62 columns]
      Var Name
```

```
#
X1    GDPSH60
X2    LIFEE060
X3      P60
X4    safrica
X5      laam
```

Implement subset selection

```
[31]: R = 0
      j_hat = None
      beta_j = None
      for xi in data.columns[1:]:
          reg_df = data.loc[:, ["gamma", xi]].dropna()
          res = sp.stats.linregress(reg_df[xi], reg_df["gamma"])
          if (res.rvalue > R):
              R = res.rvalue
              j_hat = xi
              beta_j = res.slope
      print(f"The variable chosen is {labels.loc[j_hat, 'Var Name']}. Its coefficient is {beta_j} and its R^2 is {R ** 2}")
```

The variable chosen is EQINV. Its coefficient is 0.35313560484865114 and its R^2 is 0.43268825547604167

As we can see, subset selection returns the covariate EQINV, which corresponds to equipment investment.

2. We cannot use the same trick to simplify the comparison of R^2 , since our covariates are not orthogonal. Even if we standardized our features, they would still have positive covariance.
3. We implement a greedy algorithm, where we find the best covariate, then add the next best covariate. Although this is not theoretically optimal (in fact, there is always a sequence of covariates that the greedy algorithm would add to our list, but gives the worst R^2), it is usually effective in practice.
4. Code implementation of 2.3

```
[32]: def greedy_ss(data, s = 2):
      R = 0
      j_hat = []
```

```

beta_j = []
cols = list(range(1, len(data.columns)))
num_regs = 0
while (len(j_hat) < s):
    # assume outcome is in column 0
    R_t = 0
    j_hat_t = 0
    beta_j_t = 0
    #select next best xi
    for xi in cols:
        num_regs += 1
        l = [0, xi] + j_hat
        reg_df = data.iloc[:,l].dropna()
        X = reg_df.iloc[:, 1:]
        y = reg_df.iloc[:, :1]
        reg = LinearRegression().fit(X,y)
        R2 = reg.score(X, y)
        if (R2 > R_t):
            R_t = R2
            j_hat_t = xi
            beta_j_t = reg.coef_
    if (R_t > R):
        R = R_t
        j_hat.append(j_hat_t)
        beta_j = beta_j_t
        cols.remove(j_hat_t)
    else:
        break
return (j_hat, beta_j, R, num_regs)

```

Run function

```
[33]: j_hat, beta_j, R, num_regs = greedy_ss(data, 2)
```

```
[34]: print(f"The best indices we observed with {num_regs} regressions were:\n_
      ↪{labels.loc[data.columns[j_hat], 'Var Name'].values} \nwith values\n_
      ↪{beta_j} \nand R^2: \n{R}")

```

The best indices we observed with 121 regressions were:

```
[' EQINV' ' CONFUC']
```

with values

```
[[0.08781327 0.31485877]]
```

and R^2:

```
0.5790821694783819
```

The variables we managed to select were Equipment Investment and proportion of your country confucian.

5. The first variable we selected, Equipment Investment, makes intuitive sense. That seems like

the most reliable and consistent way to encourage long term growth. However, the second variable does not make intuitive sense. There is an argument to make for the culture of Confucianism contributing to economic growth, it doesn't really mesh with our knowledge of the world that it is the second most impactful variable. I would not claim these are the variables that matter most. We would need to make the very strong assumption that there is no omitted variable bias, which would imply that the regression is causal.

6. If we were to make the very strong, and probably not true assumption, that there is no omitted variable bias with this regression, then we would be able to claim that this regression is causal.