

Homework 5

```
In [2]: #imports
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [3]: wage_raw = pd.read_csv('Wage.csv')
```

1. K-Means

1. Instead of just having an indicator for job class we add an indicator for each of the categorical variables. The categorical variables are:

- sex
- maritl
- race
- education
- region
- jobclass
- health
- health_ins

```
In [3]: cat_var = ['sex', 'maritl', 'race', 'education', 'region', 'jobclass', 'health', 'h
wage = pd.get_dummies(wage_raw, columns=cat_var, drop_first=True)
```

```
In [102... wage
```

Out[102...

	year	age	logwage	wage	maritl_2. Married	maritl_3. Widowed	maritl_4. Divorced	maritl_5. Separated	race_2. Black
0	2006	18	4.318063	75.043154	False	False	False	False	False
1	2004	24	4.255273	70.476020	False	False	False	False	False
2	2003	45	4.875061	130.982177	True	False	False	False	False
3	2003	43	5.041393	154.685293	True	False	False	False	False
4	2005	50	4.318063	75.043154	False	False	True	False	False
...
2995	2008	44	5.041393	154.685293	True	False	False	False	False
2996	2007	30	4.602060	99.689464	True	False	False	False	False
2997	2005	27	4.193125	66.229408	True	False	False	False	True
2998	2005	27	4.477121	87.981033	False	False	False	False	False
2999	2009	55	4.505150	90.481913	False	False	False	True	False

3000 rows × 18 columns



2. Code

```
In [5]: rng = np.random.default_rng()
train_idx = rng.binomial(1, 0.8, size = wage.shape[0]).astype(bool)
wage_train = wage[train_idx]
wage_test = wage[~train_idx]
```

3. We assume that everyone has completed middle school for a minimum of 9 years of schooling. If someone is a High School Grad that counts for 4 more years. If they have some college, we assume that's an associate degree. The cases where someone drops out at 1 year or 3 years average out to 2 years (most of the time). Some college should count for 2 years. College grad counts for 4 years. Advanced degrees can be masters or phds or professional degrees. Masters usually are 2 years, and Phds are usually 5-6, but often take much longer. We assume it adds a bit more years of schooling, say 4 years on average.

4. Code

```
In [104... # make education years column
education_years = 9 + wage['education_2. HS Grad'] * 4 + wage['education_3. Some Co
# make kmeans data_frame
```

```
wage_kmeans = pd.DataFrame({'jobclass': wage['jobclass_2. Information'], 'age': wag

# standardize
wage_kmeans[['age', 'education_years', 'logwage']] = (wage_kmeans[['age', 'educatio

# split
wage_kmeans_train = wage_kmeans[train_idx]
wage_kmeans_test = wage_kmeans[~train_idx]
wage_kmeans
```

Out[104...

	jobclass	age	education_years	logwage
0	False	-2.115215	-1.888180	-0.954767
1	True	-1.595392	0.577880	-1.133275
2	False	0.223986	-0.038635	0.628727
3	True	0.050712	0.577880	1.101591
4	True	0.657171	-0.655150	-0.954767
...
2995	False	0.137349	-0.038635	1.101591
2996	False	-1.075570	-0.655150	-0.147391
2997	False	-1.335481	-1.888180	-1.309956
2998	False	-1.335481	-0.038635	-0.502580
2999	False	1.090356	-0.655150	-0.422897

3000 rows × 4 columns

5. Code

In [105...

```
import sklearn.cluster
num_clusters = 8
kmeans_clf = sklearn.cluster.KMeans(n_clusters = num_clusters)
kmeans_clf.fit(wage_kmeans_train)
```

Out[105...

▼ KMeans ⓘ ?

KMeans()

6. Predict

In [106...

```
kmeans_weights = np.zeros(num_clusters)
for i in range(num_clusters):
    kmeans_weights[i] = np.mean(wage_kmeans_train.iloc[kmeans_clf.labels_ == i]['jo
print(kmeans_weights)
```

```
[0.66666667 0.50743494 0.8          0.29299363 0.46285714 0.39732143
 0.51923077 0.24096386]
```

7. Here is some code to investigate this:

```
In [107... print(kmeans_clf.cluster_centers_)

[[ 0.66812227  1.25038882  0.95478894  0.2731773 ]
 [ 0.50743494 -0.77362873 -0.03634307 -0.0313303 ]
 [ 0.8         -0.09379251  1.81091033  1.06612075]
 [ 0.29299363  0.09320234 -1.88818023 -0.5036751 ]
 [ 0.46285714  0.37102111 -0.43672755 -1.42916273]
 [ 0.39732143  0.92133658 -0.49138822 -0.02732775]
 [ 0.51910828  0.08602857  0.29514711  1.09067936]
 [ 0.24096386 -1.44717001 -0.62543847 -1.19975317]]
```

As we can see, the first coordinate of each of the cluster centers (the part corresponding to the job) is exactly the fraction we calculated prior. This makes sense since each cluster center's should just be the mean of the group, otherwise the variance would increase.

8. Code

```
In [108... kmeans_weights = kmeans_weights > 0.5
kmeans_preds_train = np.eye(num_clusters)[kmeans_clf.predict(wage_kmeans_train)] @
kmeans_err_train = np.mean(kmeans_preds_train != wage_kmeans_train['jobclass'])
print(f"The training error for the k-means classifier is {kmeans_err_train}")
```

The training error for the k-means classifier is 0.3831578947368421

9. Code

```
In [109... kmeans_preds_test = np.eye(num_clusters)[kmeans_clf.predict(wage_kmeans_test)] @ km
```

10. Code

```
In [110... kmeans_err_test = np.mean(kmeans_preds_test != wage_kmeans_test['jobclass'])
print(f"The test error for the k-means classifier is {kmeans_err_test}")
```

The test error for the k-means classifier is 0.384

2. Neural Networks

1. Code

```
In [4]: import torch
import torch.nn as nn
import torch.optim as optim
import torch.functional as F
import torch.utils.data as data
```

```
In [9]: device = torch.accelerator.current_accelerator().type if torch.accelerator.is_avail
print(f'Using device: {device}')
```

Using device: cuda

```
In [10]: class Net(nn.Module):
def __init__(self):
super(Net, self).__init__()
self.linear_relu_logit_stack = nn.Sequential(
nn.Linear(3, 16),
nn.ReLU(),
nn.Linear(16, 16),
nn.ReLU(),
nn.Linear(16, 1),
nn.Sigmoid()
)
def forward(self, x):
return self.linear_relu_logit_stack(x)
```

```
In [114... model = Net().to(device)
print(model)
```

```
Net(
  (linear_relu_logit_stack): Sequential(
    (0): Linear(in_features=3, out_features=16, bias=True)
    (1): ReLU()
    (2): Linear(in_features=16, out_features=16, bias=True)
    (3): ReLU()
    (4): Linear(in_features=16, out_features=1, bias=True)
    (5): Sigmoid()
  )
)
```

```
In [115... loss_fn = nn.BCELoss()
optimizer = optim.Adam(model.parameters(), lr=10e-4)
```

```
In [11]: wage_nn_train = torch.tensor(wage_kmeans_train[['age', 'education_years', 'logwage']
jobclass_nn_train = torch.tensor(wage_kmeans_train['jobclass'].values, dtype=torch.f
jobclass_nn_train = jobclass_nn_train.view(-1, 1)
wage_nn_test = torch.tensor(wage_kmeans_test[['age', 'education_years', 'logwage']]
jobclass_nn_test = torch.tensor(wage_kmeans_test['jobclass'].values, dtype=torch.fl
jobclass_nn_test = jobclass_nn_test.view(-1, 1)
```

```

-----
NameError                                Traceback (most recent call last)
Cell In[11], line 1
----> 1 wage_nn_train = torch.tensor(wage_kmeans_train[['age', 'education_years', 'logwage']].values, dtype=torch.float32).to(device)
      2 jobclass_nn_train = torch.tensor(wage_kmeans_train['jobclass'].values, dtype=torch.float32).to(device)
      3 jobclass_nn_train = jobclass_nn_train.view(-1, 1)

NameError: name 'wage_kmeans_train' is not defined

```

```

In [5]: def train(X,y, batch_size, model, loss_fn, optimizer):
        model.train()
        for i in np.random.permutation(range(len(X)//batch_size)):
            Xi, yi = X[i*batch_size: (i+1)*batch_size], y[i*batch_size: (i+1)*batch_size]
            Xi, yi = Xi.to(device), yi.to(device)
            # pred
            pred = model(Xi)
            loss = loss_fn(pred, yi)
            # backprop
            loss.backward()
            optimizer.step()
            optimizer.zero_grad()

            # if i % 100 == 0:
            #     loss = loss.item()
            #     print(f"Loss: {loss:>7f} [{i:>5d}/{len(X)}]")

```

```

In [6]: def test(X, y, model, loss_fn):
        model.eval()
        size = len(X)
        test_loss, correct = 0, 0
        with torch.no_grad():
            for Xi, yi in zip(X, y):
                Xi, yi = Xi.to(device), yi.to(device)
                pred = model(Xi)
                loss = loss_fn(pred, yi)
                test_loss += loss.item()
                correct += (pred.round() == yi).type(torch.float).item()
        test_loss /= size
        correct /= size
        print(f"Test Error: \n Accuracy: {(100*correct)>0.1f}%, Avg loss: {test_loss:>7f}")
        return correct, test_loss

```

```

In [6]: epochs = 5
        for t in range(epochs):
            print(f"Epoch {t+1}\n-----")
            train(wage_nn_train, jobclass_nn_train, model, loss_fn, optimizer)
            test(wage_nn_train, jobclass_nn_train, model, loss_fn)

```

Epoch 1

```

-----
NameError                                Traceback (most recent call last)
Cell In[6], line 4
      2 for t in range(epochs):
      3     print(f"Epoch {t+1}\n-----")
----> 4     train(wage_nn_train, jobclass_nn_train, model, loss_fn, optimizer)
      5     test(wage_nn_train, jobclass_nn_train, model, loss_fn)

NameError: name 'train' is not defined

```

2. Code

```

In [121... model.eval()
wage_nn_train_err = ((model(wage_nn_train) > 0.5) == jobclass_nn_train).sum().item()
print(f"The training error for the neural network is {1 - wage_nn_train_err}")

```

The training error for the neural network is 0.3701052631578947

3. Code

```

In [122... model.eval()
wage_nn_test_err = ((model(wage_nn_test) > 0.5) == jobclass_nn_test).sum().item()/1
print(f"The test error for the neural network is {1 - wage_nn_test_err}")

```

The test error for the neural network is 0.35840000000000005

3. Regression Trees

1. Code

```

In [123... wage_dt_train = wage_kmeans_train[['age', 'education_years', 'logwage']]
jobclass_dt_train = wage_kmeans_train['jobclass']
wage_dt_test = wage_kmeans_test[['age', 'education_years', 'logwage']]
jobclass_dt_test = wage_kmeans_test['jobclass']

```

```

In [124... import sklearn.tree
dt_clf = sklearn.tree.DecisionTreeClassifier(max_depth = 3)
dt_clf.fit(wage_dt_train, jobclass_dt_train)

```

```

Out[124... DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3)

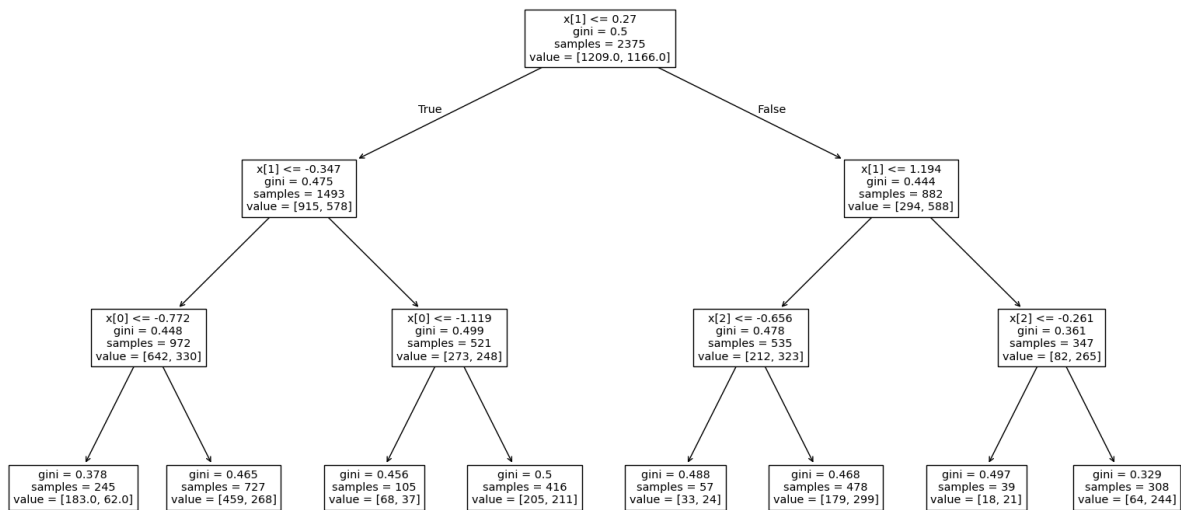
```

2. Code

```

In [125... plt.figure(figsize=(20,10))
sklearn.tree.plot_tree(dt_clf)
plt.show()

```



3. Code

```
In [126... wage_dt_train_err = 1 - dt_clf.score(wage_dt_train, jobclass_dt_train)
print(f"The training error for the decision tree is {wage_dt_train_err}")
```

The training error for the decision tree is 0.36084210526315785

4. Code

```
In [127... wage_dt_test_err = 1 - dt_clf.score(wage_dt_test, jobclass_dt_test)
print(f"The test error for the decision tree is {wage_dt_test_err}")
```

The test error for the decision tree is 0.36639999999999995

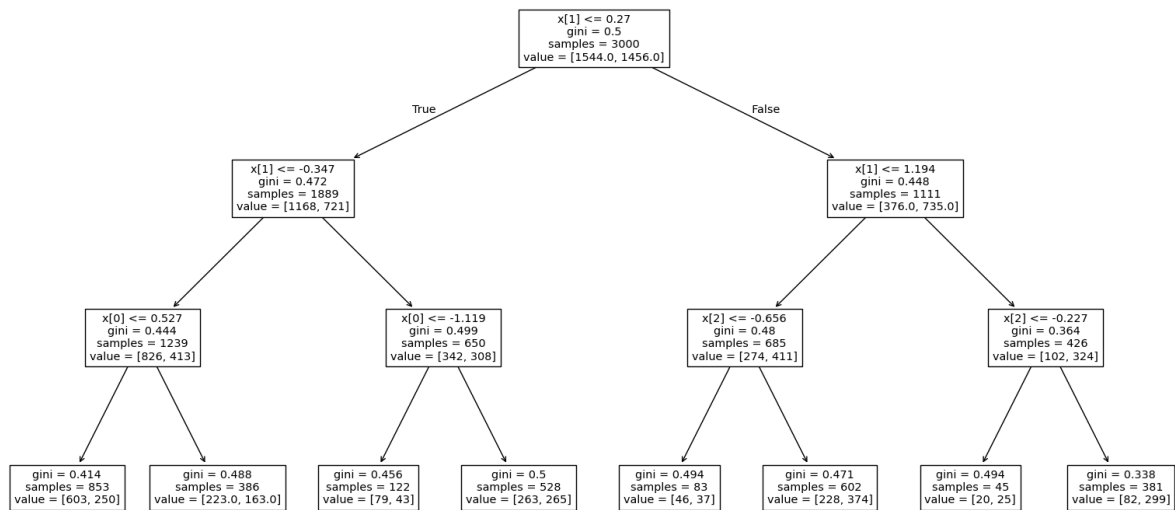
5. Code

```
In [128... wage_dt = wage_kmeans[['age', 'education_years', 'logwage']]
jobclass_dt = wage_kmeans['jobclass']
dt_full_clf = sklearn.tree.DecisionTreeClassifier(max_depth = 3)
dt_full_clf.fit(wage_dt, jobclass_dt)
```

```
Out[128... DecisionTreeClassifier
DecisionTreeClassifier(max_depth=3)
```

6. Code

```
In [129... plt.figure(figsize=(20,10))
sklearn.tree.plot_tree(dt_full_clf)
plt.show()
```

7. Code

```
In [130...] wage_dt_full_train_err = 1 - dt_full_clf.score(wage_dt, jobclass_dt)
print(f"The error on the training set for the decision tree is {wage_dt_full_train_err}")
```

The error on the training set for the decision tree is 0.362

8. Code

```
In [131...] wage_dt_full_test_err = 1 - dt_full_clf.score(wage_dt_test, jobclass_dt_test)
print(f"The error on the test set for the decision tree is {wage_dt_full_test_err}")
```

The error on the test set for the decision tree is 0.36639999999999995

4. Compare All Methods

1. Code

```
In [133...] results = pd.DataFrame({
    'kmeans': [kmeans_err_train, kmeans_err_test],
    'neural_network': [1 - wage_nn_train_err, 1 - wage_nn_test_err],
    'decision_tree': [wage_dt_train_err, wage_dt_test_err],
    'decision_tree_full': [wage_dt_full_train_err, wage_dt_full_test_err]
}, index = ['train', 'test'])
results
```

```
Out[133...]
      kmeans  neural_network  decision_tree  decision_tree_full
train  0.383158         0.370105         0.360842         0.3620
test   0.384000         0.358400         0.366400         0.3664
```

2. All of the methods performed incredibly close to each other. They were all within $3e-2$ of each other. The neural network did the best, which was expected, but all methods performed well. In theory, we should have expected the full decision tree to perform better than the standard decision tree, however, we only had 3 features, with max depth of 8, so the decision tree was quite constrained. The neural network performed the best, but not by as well as I expected. I chalk this up to there not being enough data and not using enough features.
3. (Optional) We are going to train a neural net with one more hidden layer and more intermediary nodes. We are also going to cross validate for the learning rate in the SGD optimizer and we are going to use the entire data set with the one-hot encoding.

```
In [ ]: wage = wage.astype('float64')
wage_train = wage_train.astype('float64')
wage_test = wage_test.astype('float64')
wage
```

```
In [137... class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.linear_relu_logit_stack = nn.Sequential(
            nn.Linear(17, 32),
            nn.ReLU(),
            nn.Linear(32, 32),
            nn.ReLU(),
            nn.Linear(32, 32),
            nn.ReLU(),
            nn.Linear(32, 1),
            nn.Sigmoid()
        )
    def forward(self, x):
        return self.linear_relu_logit_stack(x)
model = Net().to(device)
print(model)
```

```
Net(
  (linear_relu_logit_stack): Sequential(
    (0): Linear(in_features=17, out_features=32, bias=True)
    (1): ReLU()
    (2): Linear(in_features=32, out_features=32, bias=True)
    (3): ReLU()
    (4): Linear(in_features=32, out_features=32, bias=True)
    (5): ReLU()
    (6): Linear(in_features=32, out_features=1, bias=True)
    (7): Sigmoid()
  )
)
```

```
In [22]: wage_nn_train = torch.tensor(wage_train.drop(columns = 'jobclass_2. Information').v
jobclass_nn_train = torch.tensor(wage_train['jobclass_2. Information'].values, dtype
jobclass_nn_train = jobclass_nn_train.view(-1, 1)
```

```
wage_nn_test = torch.tensor(wage_test.drop(columns = 'jobclass_2. Information').val
jobclass_nn_test = torch.tensor(wage_test['jobclass_2. Information'].values, dtype=
jobclass_nn_test = jobclass_nn_test.view(-1, 1)
```

```
In [151... optimizer = optim.Adam(model.parameters(), lr=10e-5)
epochs = 25
```

```
In [152... for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train(wage_nn_train, jobclass_nn_train, model, loss_fn, optimizer)
    test(wage_nn_train, jobclass_nn_train, model, loss_fn)
```

Epoch 1

Test Error:

Accuracy: 57.2%, Avg loss: 0.677338

Epoch 2

Test Error:

Accuracy: 57.1%, Avg loss: 0.673603

Epoch 3

Test Error:

Accuracy: 59.4%, Avg loss: 0.667718

Epoch 4

Test Error:

Accuracy: 57.2%, Avg loss: 0.672373

Epoch 5

Test Error:

Accuracy: 59.5%, Avg loss: 0.668015

Epoch 6

Test Error:

Accuracy: 56.8%, Avg loss: 0.676829

Epoch 7

Test Error:

Accuracy: 59.7%, Avg loss: 0.668202

Epoch 8

Test Error:

Accuracy: 57.7%, Avg loss: 0.668597

Epoch 9

Test Error:

Accuracy: 57.5%, Avg loss: 0.670710

Epoch 10

Test Error:

Accuracy: 59.6%, Avg loss: 0.667262

Epoch 11

Test Error:

Accuracy: 57.4%, Avg loss: 0.673629

Epoch 12

Test Error:

Accuracy: 59.2%, Avg loss: 0.668041

Epoch 13

Test Error:

Accuracy: 59.7%, Avg loss: 0.671024

Epoch 14

Test Error:

Accuracy: 59.5%, Avg loss: 0.666076

Epoch 15

Test Error:

Accuracy: 58.1%, Avg loss: 0.668596

Epoch 16

Test Error:

Accuracy: 54.1%, Avg loss: 0.718041

Epoch 17

Test Error:

Accuracy: 58.5%, Avg loss: 0.667357

Epoch 18

Test Error:

Accuracy: 56.8%, Avg loss: 0.675292

Epoch 19

Test Error:

Accuracy: 60.0%, Avg loss: 0.665352

Epoch 20

Test Error:

Accuracy: 58.4%, Avg loss: 0.670422

Epoch 21

Test Error:

Accuracy: 56.4%, Avg loss: 0.676622

Epoch 22

Test Error:

Accuracy: 60.3%, Avg loss: 0.664770

Epoch 23

Test Error:

Accuracy: 59.6%, Avg loss: 0.664290

Epoch 24

Test Error:

Accuracy: 59.4%, Avg loss: 0.664646

Epoch 25

Test Error:

Accuracy: 56.8%, Avg loss: 0.675412

In [153...]

```
model.eval()
wage_nn_train_err = ((model(wage_nn_train) > 0.5) == jobclass_nn_train).sum().item()
print(f"The training error for the neural network is {1 - wage_nn_train_err}")
wage_nn_test_err = ((model(wage_nn_test) > 0.5) == jobclass_nn_test).sum().item()/1
print(f"The test error for the neural network is {1 - wage_nn_test_err}")
```

The training error for the neural network is 0.43200000000000005

The test error for the neural network is 0.40480000000000005

It's clear that our nn is suffering. We will drop wages because they are codependent with logwages. We also demean year and age, since only the relative age and years matter

In [7]:

```
cat_var = ['sex', 'maritl', 'race', 'education', 'region', 'jobclass', 'health', 'h
wage = pd.get_dummies(wage_raw, columns=cat_var, drop_first=True)
wage = wage.astype('float64')
wage = wage.drop(columns = 'wage')
wage['year'] = wage['year'] - wage['year'].mean()
wage['age'] = wage['age'] - wage['age'].mean()
rng = np.random.default_rng()
train_idx = rng.binomial(1, 0.8, size = wage.shape[0]).astype(bool)
wage_train = wage[train_idx]
wage_test = wage[~train_idx]
wage
```

Out[7]:

	year	age	logwage	maritl_2. Married	maritl_3. Widowed	maritl_4. Divorced	maritl_5. Separated	race_2. Black	race As
0	0.209	-24.414667	4.318063	0.0	0.0	0.0	0.0	0.0	
1	-1.791	-18.414667	4.255273	0.0	0.0	0.0	0.0	0.0	
2	-2.791	2.585333	4.875061	1.0	0.0	0.0	0.0	0.0	
3	-2.791	0.585333	5.041393	1.0	0.0	0.0	0.0	0.0	
4	-0.791	7.585333	4.318063	0.0	0.0	1.0	0.0	0.0	
...	
2995	2.209	1.585333	5.041393	1.0	0.0	0.0	0.0	0.0	
2996	1.209	-12.414667	4.602060	1.0	0.0	0.0	0.0	0.0	
2997	-0.791	-15.414667	4.193125	1.0	0.0	0.0	0.0	1.0	
2998	-0.791	-15.414667	4.477121	0.0	0.0	0.0	0.0	0.0	
2999	3.209	12.585333	4.505150	0.0	0.0	0.0	1.0	0.0	

3000 rows × 17 columns



```
In [23]: device = torch.accelerator.current_accelerator().type if torch.accelerator.is_avail
print(f'Using device: {device}')
wage_nn_train = torch.tensor(wage_train.drop(columns = 'jobclass_2. Information').v
jobclass_nn_train = torch.tensor(wage_train['jobclass_2. Information'].values, dtype
jobclass_nn_train = jobclass_nn_train.view(-1, 1)
wage_nn_test = torch.tensor(wage_test.drop(columns = 'jobclass_2. Information').val
jobclass_nn_test = torch.tensor(wage_test['jobclass_2. Information'].values, dtype=
jobclass_nn_test = jobclass_nn_test.view(-1, 1)
```

Using device: cuda

In [10]: wage_nn_train.shape

Out[10]: torch.Size([2411, 16])

```
In [24]: class Net(nn.Module):
def __init__(self, d_in = 17):
super(Net, self).__init__()
self.linear_relu_logit_stack = nn.Sequential(
nn.Linear(d_in, 32),
nn.ReLU(),
nn.Linear(32, 32),
nn.ReLU(),
nn.Linear(32, 64),
nn.ReLU(),
nn.Linear(64, 32),
nn.ReLU(),
nn.Linear(32, 16),
```

```

        nn.ReLU(),
        nn.Linear(16, 1),
        nn.Sigmoid()
    )
    def forward(self, x):
        return self.linear_relu_logit_stack(x)

```

```

In [52]: loss_fn = nn.BCELoss()
# epochs = 25
batch_size = 32
model = Net(d_in = 16).to(device)
optimizer = optim.Adam(model.parameters(), lr=10e-4)

```

```

In [53]: t = 0
accuracy = 0
accuracies = [0.5]
test_accuracies = [0.5]
while max(test_accuracies) < 0.6615384615384615:
    print(f"Epoch {t+1}\n-----")
    train(wage_nn_train, jobclass_nn_train, batch_size, model, loss_fn, optimizer)
    accuracy, loss = test(wage_nn_train, jobclass_nn_train, model, loss_fn)
    test_accuracy, test_loss = test(wage_nn_test, jobclass_nn_test, model, loss_fn)
    accuracies.append(accuracy)
    test_accuracies.append(test_accuracy)
    t += 1

```


Epoch 1

Test Error:

Accuracy: 54.0%, Avg loss: 0.688912

Test Error:

Accuracy: 51.5%, Avg loss: 0.692923

Epoch 2

Test Error:

Accuracy: 54.2%, Avg loss: 0.681890

Test Error:

Accuracy: 55.4%, Avg loss: 0.680201

Epoch 3

Test Error:

Accuracy: 62.3%, Avg loss: 0.655133

Test Error:

Accuracy: 60.3%, Avg loss: 0.660713

Epoch 4

Test Error:

Accuracy: 63.3%, Avg loss: 0.647987

Test Error:

Accuracy: 62.4%, Avg loss: 0.647018

Epoch 5

Test Error:

Accuracy: 63.9%, Avg loss: 0.641303

Test Error:

Accuracy: 63.4%, Avg loss: 0.637046

Epoch 6

Test Error:

Accuracy: 63.6%, Avg loss: 0.638212

Test Error:

Accuracy: 63.1%, Avg loss: 0.633663

Epoch 7

Test Error:

Accuracy: 64.1%, Avg loss: 0.638521

Test Error:

Accuracy: 63.4%, Avg loss: 0.638104

Epoch 8

Test Error:

Accuracy: 64.0%, Avg loss: 0.634455

Test Error:

Accuracy: 65.0%, Avg loss: 0.625823

Epoch 9

Test Error:

Accuracy: 64.7%, Avg loss: 0.633673

Test Error:

Accuracy: 63.1%, Avg loss: 0.628517

Epoch 10

Test Error:

Accuracy: 64.9%, Avg loss: 0.633447

Test Error:

Accuracy: 64.4%, Avg loss: 0.626766

Epoch 11

Test Error:

Accuracy: 64.4%, Avg loss: 0.632701

Test Error:

Accuracy: 64.3%, Avg loss: 0.626422

Epoch 12

Test Error:

Accuracy: 64.7%, Avg loss: 0.630944

Test Error:

Accuracy: 64.6%, Avg loss: 0.623650

Epoch 13

Test Error:

Accuracy: 64.5%, Avg loss: 0.631863

Test Error:

Accuracy: 63.6%, Avg loss: 0.629885

Epoch 14

Test Error:

Accuracy: 63.4%, Avg loss: 0.638376

Test Error:

Accuracy: 62.7%, Avg loss: 0.635124

Epoch 15

Test Error:

Accuracy: 64.8%, Avg loss: 0.627940

Test Error:

Accuracy: 62.9%, Avg loss: 0.630257

Epoch 16

Test Error:

Accuracy: 65.3%, Avg loss: 0.628684

Test Error:

Accuracy: 62.9%, Avg loss: 0.631827

Epoch 17

Test Error:

Accuracy: 65.6%, Avg loss: 0.624737

Test Error:

Accuracy: 65.1%, Avg loss: 0.624063

Epoch 18

Test Error:

Accuracy: 65.6%, Avg loss: 0.625201

Test Error:

Accuracy: 63.2%, Avg loss: 0.627675

Epoch 19

Test Error:

Accuracy: 65.3%, Avg loss: 0.624729

Test Error:

Accuracy: 64.8%, Avg loss: 0.626835

Epoch 20

Test Error:

Accuracy: 65.7%, Avg loss: 0.625187

Test Error:

Accuracy: 63.4%, Avg loss: 0.629846

Epoch 21

Test Error:

Accuracy: 66.3%, Avg loss: 0.620099

Test Error:

Accuracy: 63.4%, Avg loss: 0.627794

Epoch 22

Test Error:

Accuracy: 65.8%, Avg loss: 0.619355

Test Error:

Accuracy: 64.4%, Avg loss: 0.627411

Epoch 23

Test Error:

Accuracy: 66.0%, Avg loss: 0.618179

Test Error:

Accuracy: 62.4%, Avg loss: 0.629185

Epoch 24

Test Error:

Accuracy: 66.4%, Avg loss: 0.616349

Test Error:

Accuracy: 63.1%, Avg loss: 0.629715

Epoch 25

Test Error:

Accuracy: 66.3%, Avg loss: 0.618377

Test Error:

Accuracy: 64.6%, Avg loss: 0.627575

Epoch 26

Test Error:

Accuracy: 66.5%, Avg loss: 0.615430

Test Error:

Accuracy: 63.4%, Avg loss: 0.623028

Epoch 27

Test Error:

Accuracy: 65.3%, Avg loss: 0.621508

Test Error:

Accuracy: 62.9%, Avg loss: 0.631204

Epoch 28

Test Error:

Accuracy: 66.2%, Avg loss: 0.613335

Test Error:

Accuracy: 64.3%, Avg loss: 0.626733

Epoch 29

Test Error:

Accuracy: 65.5%, Avg loss: 0.620411

Test Error:

Accuracy: 63.8%, Avg loss: 0.632448

Epoch 30

Test Error:

Accuracy: 66.7%, Avg loss: 0.613138

Test Error:

Accuracy: 63.8%, Avg loss: 0.627666

Epoch 31

Test Error:

Accuracy: 66.5%, Avg loss: 0.612773

Test Error:

Accuracy: 61.5%, Avg loss: 0.637041

Epoch 32

Test Error:

Accuracy: 66.8%, Avg loss: 0.607241

Test Error:

Accuracy: 62.9%, Avg loss: 0.627751

Epoch 33

Test Error:

Accuracy: 66.7%, Avg loss: 0.606945

Test Error:

Accuracy: 64.1%, Avg loss: 0.629824

Epoch 34

Test Error:

Accuracy: 66.5%, Avg loss: 0.605591

Test Error:

Accuracy: 65.5%, Avg loss: 0.627176

Epoch 35

Test Error:

Accuracy: 66.6%, Avg loss: 0.610121

Test Error:

Accuracy: 60.9%, Avg loss: 0.636251

Epoch 36

Test Error:

Accuracy: 67.1%, Avg loss: 0.604154

Test Error:

Accuracy: 61.9%, Avg loss: 0.635524

Epoch 37

Test Error:

Accuracy: 67.1%, Avg loss: 0.608591

Test Error:

Accuracy: 61.9%, Avg loss: 0.633077

Epoch 38

Test Error:

Accuracy: 67.7%, Avg loss: 0.602258

Test Error:

Accuracy: 62.6%, Avg loss: 0.626677

Epoch 39

Test Error:

Accuracy: 66.6%, Avg loss: 0.604702

Test Error:

Accuracy: 62.1%, Avg loss: 0.631172

Epoch 40

Test Error:

Accuracy: 66.7%, Avg loss: 0.607520

Test Error:

Accuracy: 61.9%, Avg loss: 0.635518

Epoch 41

Test Error:

Accuracy: 66.8%, Avg loss: 0.602754

Test Error:

Accuracy: 62.6%, Avg loss: 0.636054

Epoch 42

Test Error:

Accuracy: 67.4%, Avg loss: 0.601970

Test Error:

Accuracy: 62.6%, Avg loss: 0.643882

Epoch 43

Test Error:

Accuracy: 67.8%, Avg loss: 0.596927

Test Error:

Accuracy: 63.4%, Avg loss: 0.635919

Epoch 44

Test Error:

Accuracy: 67.9%, Avg loss: 0.594442

Test Error:

Accuracy: 63.6%, Avg loss: 0.632562

Epoch 45

Test Error:

Accuracy: 67.7%, Avg loss: 0.595256

Test Error:

Accuracy: 65.5%, Avg loss: 0.624834

Epoch 46

Test Error:

Accuracy: 67.8%, Avg loss: 0.592588

Test Error:

Accuracy: 62.9%, Avg loss: 0.630867

Epoch 47

Test Error:

Accuracy: 67.4%, Avg loss: 0.593409

Test Error:

Accuracy: 64.8%, Avg loss: 0.635298

Epoch 48

Test Error:

Accuracy: 68.6%, Avg loss: 0.594090

Test Error:

Accuracy: 63.8%, Avg loss: 0.637697

Epoch 49

Test Error:

Accuracy: 68.6%, Avg loss: 0.598494

Test Error:

Accuracy: 63.6%, Avg loss: 0.634642

Epoch 50

Test Error:

Accuracy: 67.5%, Avg loss: 0.592970

Test Error:

Accuracy: 62.9%, Avg loss: 0.634516

Epoch 51

Test Error:

Accuracy: 68.7%, Avg loss: 0.584959

Test Error:

Accuracy: 64.4%, Avg loss: 0.639305

Epoch 52

Test Error:

Accuracy: 67.2%, Avg loss: 0.600400

Test Error:

Accuracy: 61.4%, Avg loss: 0.643227

Epoch 53

Test Error:

Accuracy: 69.1%, Avg loss: 0.592140

Test Error:

Accuracy: 62.9%, Avg loss: 0.640318

Epoch 54

Test Error:

Accuracy: 69.2%, Avg loss: 0.583685

Test Error:

Accuracy: 63.1%, Avg loss: 0.635839

Epoch 55

Test Error:

Accuracy: 69.5%, Avg loss: 0.582836

Test Error:

Accuracy: 61.7%, Avg loss: 0.635465

Epoch 56

Test Error:

Accuracy: 69.4%, Avg loss: 0.578892

Test Error:

Accuracy: 65.6%, Avg loss: 0.634814

Epoch 57

Test Error:

Accuracy: 70.4%, Avg loss: 0.577554

Test Error:

Accuracy: 63.2%, Avg loss: 0.644747

Epoch 58

Test Error:

Accuracy: 70.6%, Avg loss: 0.575099

Test Error:

Accuracy: 62.9%, Avg loss: 0.642703

Epoch 59

Test Error:

Accuracy: 69.3%, Avg loss: 0.578011

Test Error:

Accuracy: 62.6%, Avg loss: 0.653849

Epoch 60

Test Error:

Accuracy: 70.1%, Avg loss: 0.571171

Test Error:

Accuracy: 61.7%, Avg loss: 0.646427

Epoch 61

Test Error:

Accuracy: 70.5%, Avg loss: 0.572552

Test Error:

Accuracy: 63.9%, Avg loss: 0.652813

Epoch 62

Test Error:

Accuracy: 70.5%, Avg loss: 0.565000

Test Error:

Accuracy: 63.8%, Avg loss: 0.647250

Epoch 63

Test Error:

Accuracy: 69.7%, Avg loss: 0.576659

Test Error:

Accuracy: 63.2%, Avg loss: 0.664534

Epoch 64

Test Error:

Accuracy: 71.1%, Avg loss: 0.569392

Test Error:

Accuracy: 64.1%, Avg loss: 0.641037

Epoch 65

Test Error:

Accuracy: 70.2%, Avg loss: 0.571338

Test Error:

Accuracy: 63.8%, Avg loss: 0.658674

Epoch 66

Test Error:

Accuracy: 69.2%, Avg loss: 0.568884

Test Error:

Accuracy: 59.7%, Avg loss: 0.667272

Epoch 67

Test Error:

Accuracy: 71.3%, Avg loss: 0.562705

Test Error:

Accuracy: 61.7%, Avg loss: 0.662775

Epoch 68

Test Error:

Accuracy: 71.7%, Avg loss: 0.557770

Test Error:

Accuracy: 64.1%, Avg loss: 0.659682

Epoch 69

Test Error:

Accuracy: 71.3%, Avg loss: 0.556723

Test Error:

Accuracy: 62.6%, Avg loss: 0.668302

Epoch 70

Test Error:

Accuracy: 69.9%, Avg loss: 0.567010

Test Error:

Accuracy: 63.1%, Avg loss: 0.668768

Epoch 71

Test Error:

Accuracy: 72.6%, Avg loss: 0.553586

Test Error:

Accuracy: 63.1%, Avg loss: 0.658022

Epoch 72

Test Error:

Accuracy: 73.0%, Avg loss: 0.550672

Test Error:

Accuracy: 64.3%, Avg loss: 0.660483

Epoch 73

Test Error:

Accuracy: 71.3%, Avg loss: 0.558513

Test Error:

Accuracy: 64.3%, Avg loss: 0.682630

Epoch 74

Test Error:

Accuracy: 73.7%, Avg loss: 0.541067

Test Error:

Accuracy: 63.4%, Avg loss: 0.659616

Epoch 75

Test Error:

Accuracy: 73.6%, Avg loss: 0.550955

Test Error:

Accuracy: 63.9%, Avg loss: 0.665608

Epoch 76

Test Error:


Accuracy: 73.7%, Avg loss: 0.538043

Test Error:

Accuracy: 66.3%, Avg loss: 0.664403

```
In [54]: import os
os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
plt.figure()
plt.plot(accuracies, label = 'train')
plt.plot(test_accuracies, label = 'test')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
```

```
plt.title('Training Accuracy')
plt.legend()
plt.show()
print(np.argmax(test_accuracies))
print(np.max(test_accuracies))
```



76

0.6632478632478632

```
In [60]: wage_nn_train = wage_nn_train.to(device)
jobclass_nn_train = jobclass_nn_train.to(device)
wage_nn_test = wage_nn_test.to(device)
jobclass_nn_test = jobclass_nn_test.to(device)
wage_nn_train_err = ((model(wage_nn_train.to(device)) > 0.5) == jobclass_nn_train).
print(f"The training error for the neural network is {1 - wage_nn_train_err}")
wage_nn_test_err = ((model(wage_nn_test) > 0.5) == jobclass_nn_test).sum().item()/1
print(f"The test error for the neural network is {1 - wage_nn_test_err}")
```

The training error for the neural network is 0.26252587991718423

The test error for the neural network is 0.3367521367521368

This test error beats our other test error by 3%!